

Sentiment Analysis on IMDB Movie Reviews

Evaluation and comparison of multi machine learning approaches to predict a binary sentiment of a movie review

Roman Studer

BSc Data Science Student,
FHNW Brugg-Windisch, Switzerland
roman.studer1@students.fhnw.ch

Lukas Gehrig

BSc Data Science Student,
FHNW Brugg-Windisch, Switzerland
lukas.gehrig@students.fhnw.ch

Vincenzo Timmel

BSc Data Science Student,
FHNW Brugg-Windisch, Switzerland
vincenzo.timmel@students.fhnw.ch

Abstract— In this report we look at the binary sentiment classification of reviews (Positive or Negative sentiment) with different Machine Learning algorithms. From classical algorithms like linear regression to deep learning models. We explain how the data was prepared and how to interpret the results of the models.

Keywords—component, formatting, style, styling, insert (key words)

I. INTRODUCTION

A. Goal of the analysis

This text classification task is focused on the evaluation and comparison of several natural language processing methods for sentiment classification. For this purpose, we use a dataset with 50 thousand film reviews from the IMDB¹ platform. A dataset that is often used for benchmarks in the NLP environment. The dataset is suitable for a binary sentiment classification as it is labeled with the binary feature "sentiment" which either takes on the value "positive", if the sentiment of the review is positive, or "negative" if the sentiment is negative [1]. An example record of the dataset would be:

Table 1, Example of a review in the IMDB dataset

Index	Review	Sentiment
124	"I like this movie"	positive

B. Process

At the start of the process, functions were defined that allow a uniform pre-processing of the data. This ensures that all employed models receive the same data as input. In chapter II we describe the pre-processing in more detail. With the pre-processed data several models were trained and their accuracy on train and test set was recorded. In a further step, the models have been compared. Assumptions about the results, how they were obtained and whether they are robust, are given at the end of the report in chapter 5.

II. PREPROCESSING

Language offers us humans a way to communicate complex issues in an adaptable and national, even universal, way. This complexity is reflected in the almost endless possibilities of language use. Since machine learning models require numerical values as input, a transformation of the text must

take place. In most cases, the individual words of a sentence are divided into a list and each of them is transformed into its basic form. The process of transformation into the basic form is called stemming. Stemming (stem form reduction, normal form reduction) is the term used in information retrieval as well as in linguistic computer science to describe a procedure by which different morphological variants of a word are reduced to their common root, e.g., the declension of "Wortes" to "Wort" and conjugation of "gesehen" or "sah" to "seh". Models like BERT or T5 can also use unprocessed texts as input [2]. The following steps were performed on the IMDB Movie Review dataset:

1. Removing the HTML tags. These are not relevant for understanding the input.
2. removing brackets with text (it seems that bracket text was a hyperlink or quote or reference).
3. removing special characters, e. g. ?, !, /
4. removing stop words. Stop words are common or "filler" words, which contain little information. For example, connection words. They are removed.
5. put words into the basic form with the Porter-Stemmer-Algorithm, which applies multiple, hardcoded rules to reduce the word-length. Some examples: `likes`, `liked`, `likely` and `liking` will all be reduced to `like`.

For Transformer models like BERT, steps 4. and 5. are not applied. The target variable Sentiment was also adjusted. An encoding of the binary target variable was made. A positive review is now annotated with the integer 1, a negative review with the integer 0. Table 2 shows an example of the prepared text:

Table 2, Example of a review after preprocessing

Index	Review	Sentiment
15594	["saw", "thi", "movi", "wa",..]	0

A. Train-Test-Split

For reproducibility, a function was provided which returns the same split of input data for all models. The 50 thousand reviews are split 80% (40 thousand reviews) into a test set. The training set consists of the remaining 20% (10 thousand reviews).

¹ IMDB: Internet Movie Database

III. MODELLING

A. Baseline Model

As a basic comparison, a baseline model was created. The approach of this model is to count positive and negative words in a review and suggest a positive or negative sentiment based on this count. A list of positive words as well as a list of negative words was imported. Since there is no learning algorithm behind this method, the score was calculated on the test dataset only. For this method, both the list of negative and positive words and the IMDB-Dataset was pre-processed with all the steps mentioned in Chapter 2.

B. Evaluation metric

Because the dataset is perfectly balanced, we use the metric Accuracy and a confusion matrix, to see if our model maybe handles one class much better than the other.

C. Matrix representation

1) Bag of Words

Bag of Words, a method often used for document classification based on the frequency of terms, converts a collection of texts/documents into a matrix with the token frequency as elements. Where the row represents the respective review, and the column represents a word. Grammar and arrangement of input is not included in this model. In the trainset, a total of 156'200 unique words were counted after pre-processing. Bag of Words was used as input of a Logistic Regression, SVM as well as Multinomial Naive Bayes.

2) TF-IDF

TF-IDF stands for Term Frequency - Inverse Document Frequency and is calculated for each word in the corpus. This value gives an insight into the importance of a term. The Term Frequency, i.e., the relative frequency of a term in a document combined with the relative, logarithmic frequency over all documents (formula 1) gives words that occur frequently and, in many documents, less weight than words that only occur in a few documents.

$$tf(t, d) = \frac{f_{t,d}}{\sum_{t' \in d} f_{t',d}}$$

$$idf(t, D) = \log \frac{N}{|d \in D : t \in d|}$$

Where TF-IDF is a combination of the two formulas:

$$tf - idf(t, d, D) = tf(t, d) \cdot idf(t, D)$$

The result of TF-IDF, when using Sklearns is a matrix with the dimensions $(n_samples, 156'136)$, when only *unigrams* are used. This is very close to the number of unique words. It is not exactly the same, because **sklearn.feature_extraction.text.TfidfVectorizer/CountVectorizer** uses a regex pattern, which was able to extract words attached or surrounded by either // or ^, while our method

was not. In future, our method to remove special characters should be improved.

D. Classic models

In addition to the baseline model, classical machine learning models were used at the beginning. These included Logistic Regression with l2 regularization, support vector machines with hinge loss and a multinomial Naive Bayes model. Accuracy was measured for the three models using Bag of Words and TF-IDF as input. The best regularization-factor for each model was found via a 5-fold cross validation grid search, using **sklearn.model_selection.GridSearchCV**.

E. Deepl Learning Models

As an extension, two deep learning models were trained. The data used were prepared with the same preprocessing steps as described in chapter II..

1) Bidirectional LSTM

LSTM uses a pre-trained word embedding. Due to memory constraints, the fastText Mini version was used for this case. GloVe (Global Vectors), another word embedding, was also imported for comparison. Thus, two models with different embeddings were trained. The LSTM model was trained with a maximum length of 128, a training batch size of 16, validation batch size of 6 and with 5 epochs.

2) BERT

Another language model, BERT, was trained on a corpus with 800 million words from BooksCorpus and 2,500 million words from Wikipedia. The input data was prepared as in Chapter II, except for steps 4 and 5. A maximum length of 512, a batch size of 16 was used for 4 epochs.

Accuracy was recorded for all models on training and test set. The results are discussed in the following chapter.

IV. EVALUATION

The evaluation of the models was performed by the metric Accuracy. Accuracy indicates the relative frequency of a correct prediction. A well-known metric for a binary classification. For learning algorithms such as bidirectional LSTMs or BERT, in addition to the final score, the progression of the metric over the training epochs is interesting and important for detecting overfitting.

Model	(min_df, max_df)	N_gram Range, Type	Accuracy	
			Train	Test
Baseline	-	-	-	0.7131
LR (TF-IDF)	(0, 1)	(1,3), word	0.99628	0.6061
LR (BoW)	(0, 1)	(1,3), word	0.9963	0.715
SVM (TF-IDF)	(0, 1)	(1,3), word	0.99628	0.7485
SVM (BoW)	(0, 1)	(1,3), word	0.99626	0.5091
MNB (TF-IDF)	(0, 1)	(1,3), word	0.99626	0.7512
MNB (BoW)	(0, 1)	(1,3), word	0.99985	0.8857

LR (TF-IDF)	(0.001, 0.999)	(1,3), word	0.9294	0.8986
LR (BoW)	(0.001, 0.999)	(1,3), word	0.9645	0.8925
LR (TF-IDF)	(0.001, 0.999)	(3,3), word	0.6702	0.6569
LR (BoW)	(0.001, 0.999)	(3,3), word	0.66968	0.6552
SVM (TF-IDF)	(0.001, 0.999)	(3,3), word	0.9313	0.899
SVM (BoW)	(0.001, 0.999)	(3,3), word	0.9671	0.8938
SVM (TF-IDF)	(0.1, 0.9)	(3,3), word	0.768	0.7632
SVM (BoW)	(0.1, 0.9)	(3,3), word	0.7678	0.7625
SVM (TF-IDF)	(0.001, 0.999)	(1,4), word	0.93145	0.8991
SVM (BoW)	(0.001, 0.999)	(1,4), word	0.96698	0.8936
SVM (TF-IDF)	(0.001, 0.999)	(1,4), char	0.9112	0.8732
SVM (BoW)	(0.001, 0.999)	(1,4), char	0.92257	0.8647
SVM (TF-IDF)	(0.001, 0.999)	(1,3), char, word	0.93887	0.9013
MNB (TF-IDF)	(0.001, 0.999)	(1,3), word	0.89095	0.8727
MNB (BoW)	(0.001, 0.999)	(1,3), word	0.8757	0.8642
MNB (TF-IDF)	(0.001, 0.999)	(1,4), char	0.82485	0.8142
MNB (BoW)	(0.001, 0.999)	(1,4), char	0.78382	0.7775
LSTM (fastText)	-	-	0.99055	0.8358
LSTM (GloVe)	-	-	0.9355	0.8739
BERT	-	-	0.87663	0.7186

We see a high score for the training set and a significantly lower score for the test set in all models. This could indicate

overfitting of the models. Especially for the BERT model, no significant improvement to the baseline was achieved. This is partly due to the fact that the model actually overfitted. Training with a higher batch size, in addition to a high dropout rate could mitigate this. However, due to hardware constraints, this option could not be tested. SVC with both word and char n_grams with a range of (1,3) provides by far the best result with an improvement of over 19% over the baseline.

SVM Multinomial Naïve Bayes and logistic regression all provide a very good result. However, at the beginning, there was an error with the parameter *min_df* and *max_df* in **Sklearn.feature_extraction.text**; from the notebook [3] we based part of this mini challenge on, the values for *min_df*/*max_df* were set on 0 and 1, respectively. Because we ignore all words which appear at most in 1 document, we are only able to overfit on the training-data. However, when using meaningful parameter, all LR, SVM and MNB achieve good results. There also seems to be combinations of Models and Word-Vectorizing methods, which are very robust to the mistake with *min_df*/*max_df*, which could lead to some interesting findings.

REFERENCES

- [1] "Sentiment Analysis." <http://ai.stanford.edu/~amaas/data/sentiment/> (accessed Oct. 31, 2021).
- [2] "nlp - Using trained BERT Model and Data Preprocessing," *Stack Overflow*. <https://stackoverflow.com/questions/63979544/using-trained-bert-model-and-data-preprocessing> (accessed Nov. 03, 2021).
- [3] "Sentiment Analysis of IMDB Movie Reviews" *kaggle* <https://www.kaggle.com/lakshmi25npathi/sentiment-analysis-of-imdb-movie-reviews> (accessed Nov. 09, 2021).

Tables

- Table 1, Example of a review in the IMDB dataset 1
Table 2, Example of a review after preprocessing 1